

R Tutorial

Hai Yang

October 25, 2016

R

R is 'GNU S', a freely available language and environment for statistical computing and graphics which provides a wide variety of statistical and graphical techniques

- ▶ linear and nonlinear modeling
- ▶ statistical tests
- ▶ time series analysis
- ▶ classification
- ▶ clustering
- ▶ genomics

Download and install R (precompiled binary distribution) from the nearest R mirror site, <https://cran.stat.ucla.edu>

Look and feel of R

```
1+2
```

```
## [1] 3
```

```
2*2
```

```
## [1] 4
```

```
2*2*2*2
```

```
## [1] 16
```

```
2^4
```

```
## [1] 16
```

More calculations in R

```
log(16, 2)
```

```
## [1] 4
```

```
pi
```

```
## [1] 3.141593
```

```
pi*(3^2)
```

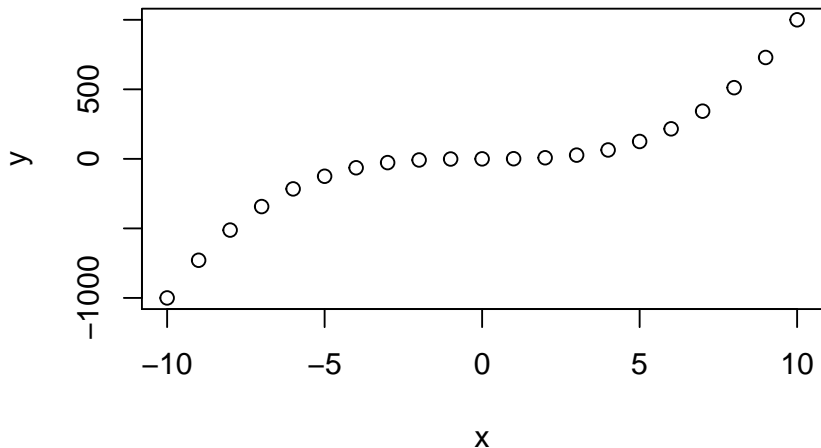
```
## [1] 28.27433
```

```
sin(pi/2)
```

```
## [1] 1
```

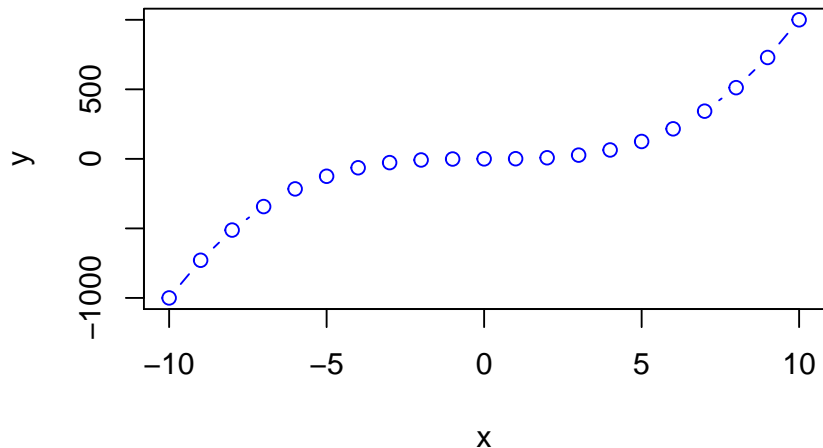
Simple plot in R

```
x = c(-10:10)
y = x^3
plot(x, y)
```



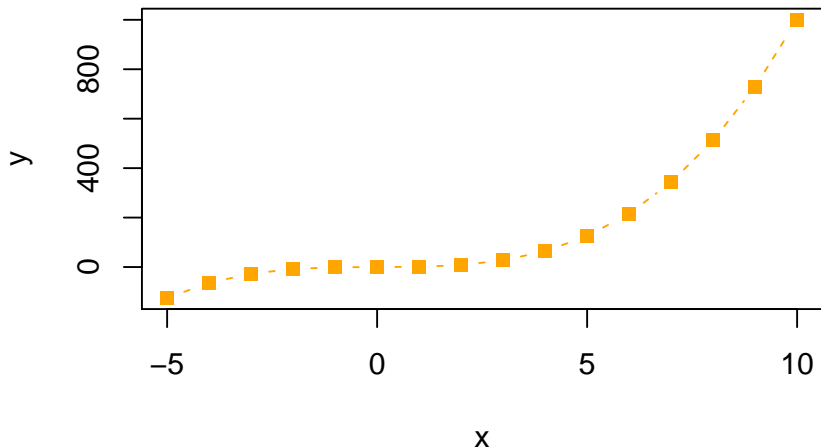
Make the figure fancier by adding more plotting parameters

```
x = c(-10:10)
y = x^3
plot(x, y, pch=1, lty=1, type="b", col="blue")
```



Try changing parameter values

```
x = c(-5:10)
y = x^3
plot(x, y, pch=15, lty=2, type="b", col="orange")
```



R Studio

RStudio is a set of integrated tools designed to help you be more productive with R. It includes

- ▶ console
- ▶ syntax-highlighting editor that supports direct code execution
- ▶ plotting
- ▶ history
- ▶ debugging
- ▶ workspace management.

Download and install RStudio from

<https://www.rstudio.com/products/rstudio/download>

R Studio text editor

R Studio -> File -> New File -> R Script to create an R code file, type commands below in the text file, click the first line, and click Run button three times.

```
x = c(-10:10)
y = x^3
plot(x, y, pch=1, lty=1, type="b", col="blue")
```

R Packages

R packages are collections of functions, data, and compiled code in a ready-to-use format with enriched documentation.

Use R Studio to install R packages R Studio -> Tools -> Install Packages... -> type in package name like "ggplot2"

- ▶ package name 1 "xlsx"
- ▶ package name 2 "foreign"
- ▶ package name 3 "ggplot2"

R Objects

R basic command

- ▶ Data input
- ▶ Data Management
- ▶ Basic Statistics
- ▶ Basic Graphs

Data input

Data input:

- ▶ Importing Data
- ▶ Data export
- ▶ Viewing Data
- ▶ Missing Data

Importing Data

```
## Keyboard input
# create a data frame from scratch
age <- c(25, 30, 56)
gender <- c("male", "female", "male")
weight <- c(160, 110, 220)
mydata <- data.frame(age,gender,weight)

## From web
# mydata <- read.csv("http://www.ats.ucla.edu/stat/data/hd")

## From text file
# mydata <- read.table("/Users/hai/Dropbox (DBMI)/Rtutorial")

## From Excel
library(xlsx)

## Loading required package: rJava
```

Data export

To A Tab Delimited Text File

```
write.table(mydata, "/Users/hai/Dropbox (DBMI)/Rtutorial/my
```

To an Excel Spreadsheet

```
library(xlsx)
```

```
# write.xlsx(mydata, "/Users/hai/Dropbox (DBMI)/Rtutorial/
```

Viewing Data

```
# list objects in the working environment  
ls()
```

```
## [1] "age"      "gender" "mydata" "weight" "x"      "y"
```

```
# list the variables in mydata  
names(mydata)
```

```
## [1] "Pat_id" "x"      "y"
```

```
# list the structure of mydata  
str(mydata)
```

```
## 'data.frame':    20 obs. of  3 variables:  
## $ Pat_id: num  1 2 3 4 5 6 7 8 9 10 ...  
## $ x      : num  7163 6443 6801 6236 3998 ...  
## $ y      : num  5.27 4.98 5.2 4.87 4.65 ...
```


Missing Data

```
# Testing for Missing Values
```

```
x <- c(1,2,NA,3)
```

```
is.na(x)
```

```
## [1] FALSE FALSE  TRUE FALSE
```

```
y <- c(1,2,3,NA)
```

```
is.na(y)
```

```
## [1] FALSE FALSE FALSE  TRUE
```

```
# Excluding Missing Values from Analyses
```

```
mean(x) # returns NA
```

```
## [1] NA
```

```
mean(x, na.rm=TRUE) # returns 2
```

Data Management

Data Management:

- ▶ Creating new variables
- ▶ Control Structures
- ▶ Sorting Data and Merging Data
- ▶ Subsetting Data
- ▶ Using `with()` and `by()`

Creating new variables

```
# Creating new variables
```

```
mydata$x <- c(1,2,NA,3)
```

```
mydata$y <- c(1,2,3,NA)
```

```
mydata$sum <- mydata$x + mydata$y
```

```
mydata$mean <- (mydata$x + mydata$y)/2
```

```
# Recoding variables
```

```
# create 2 age categories
```

```
mydata$age <- c(37,49,79,59,49,90,37,49,79,59,49,90,37,49,7
```

```
mydata$agecat <- ifelse(mydata$age > 70,
```

```
c("older"), c("younger"))
```

Control Structures

- ▶ if-else
if (cond) expr
if (cond) expr1 else expr2
- ▶ for
for (var in seq) expr
- ▶ while
while (cond) expr
- ▶ switch
switch(expr, ...)
- ▶ ifelse
ifelse(test,yes,no)

Sorting Data and Merging Data

```
# sort by mpg and cyl
attach(mtcars)
newdata <- mtcars[order(mtcars$mpg, mtcars$cyl),]

# Adding Columns
# merge two data frames by ID and Country
mtcars$ID <- rownames(mtcars)
data_frameA <- mtcars[,c(1,2,3,4,5,6,12)]
data_frameB <- mtcars[,c(6,7,8,9,10,11,12)]
total <- merge(data_frameA,data_frameB,by=c("ID","wt"))

# Adding Rows
data_frameA <- head(mtcars)
data_frameB <- tail(mtcars)
total <- rbind(data_frameA, data_frameB)
```

Subsetting Data part 1

Selecting (Keeping) Variables

#select by variable names "mpg", "cyl", "disp"

```
mydata <- mtcars
```

```
myvars <- c("mpg", "cyl", "disp")
```

```
newdata <- mydata[myvars]
```

select 1st and 5th thru 10th variables

```
newdata <- mydata[c(1,5:10)]
```

Excluding (DROPPING) Variables

#exclude variables "mpg", "cyl", "disp"

```
mydata <- mtcars
```

```
myvars <- names(mydata) %in% c("mpg", "cyl", "disp")
```

```
newdata <- mydata[!myvars]
```

exclude 3rd and 5th variable

```
newdata <- mydata[c(-3,-5)]
```

Subsetting Data part 1

Selecting Observations

#first 5 observations

```
newdata <- mydata[1:5,]
```

based on variable values

```
newdata <- mydata[ which(gear== 3 & wt > 65),]
```

Selection using the Subset Function

#using subset function

```
newdata <- subset(mydata, mpg >= 20 | mpg < 30,  
select=c(ID, wt))
```

Using with() and by()

- ▶ With

`with(data, expression)`

example applying a t-test to a data frame mydata

`with(mydata, t.test(y ~ group))`

- ▶ By

`by(data, factorlist, function)`

example obtain variable means separately for
each level of byvar in data frame mydata

`by(mydata, mydata$byvar, function(x) mean(x))`

Basice Statistics

Basice Statistics

- t-test & Correlations
- ANOVA
- Multiple Regression
- Regression Diagnostics

t-test & Correlations

► Correlations

```
# Correlation matrix from mtcars  
# with mpg, cyl, and disp as rows  
# and hp, drat, and wt as columns  
x <- mtcars[1:3]  
y <- mtcars[4:6]  
cor(x, y)
```

```
##           hp      drat      wt  
## mpg  -0.7761684  0.6811719 -0.8676594  
## cyl   0.8324475 -0.6999381  0.7824958  
## disp  0.7909486 -0.7102139  0.8879799
```

► t-tests

```
# independent 2-group t-test  
t.test(mtcars$mpg, mtcars$wt)
```

ANOVA

```
# Analysis of Covariance
```

```
fit <- aov(mpg ~ wt + cyl, data=mtcars)
```

```
fit
```

```
## Call:
```

```
##      aov(formula = mpg ~ wt + cyl, data = mtcars)
```

```
##
```

```
## Terms:
```

```
##                wt                cyl Residuals
```

```
## Sum of Squares  847.7252   87.1500   191.1720
```

```
## Deg. of Freedom      1           1         29
```

```
##
```

```
## Residual standard error: 2.567516
```

```
## Estimated effects may be unbalanced
```

Multiple Regression

```
## Fitting the Model
```

```
# Multiple Linear Regression Example
```

```
fit <- lm(mpg ~ wt + cyl + hp, data=mtcars)
```

```
summary(fit) # show results
```

```
##
```

```
## Call:
```

```
## lm(formula = mpg ~ wt + cyl + hp, data = mtcars)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -3.9290 -1.5598 -0.5311  1.1850  5.8986
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept) 38.75179      1.78686  21.687  < 2e-16 ***
```

```
## wt          -3.16697      0.74058  -4.276 0.000199 ***
```

Multiple Regression (cont)

```
## Variable Selection
#Stepwise Regression
library(MASS)
fit <- lm(mpg ~ wt + cyl + hp, data=mtcars)
step <- stepAIC(fit, direction="both")
```

```
## Start:  AIC=62.66
## mpg ~ wt + cyl + hp
##
##           Df Sum of Sq    RSS    AIC
## <none>                176.62 62.665
## - hp      1      14.551 191.17 63.198
## - cyl     1      18.427 195.05 63.840
## - wt      1     115.354 291.98 76.750
```

```
step$anova # display results
```

```
## Stepwise Model Path
```

Regression Diagnostics

```
## Assume that we are fitting a multiple linear regression  
  
library(car)  
fit <- lm(mpg~disp+hp+wt+drat, data=mtcars)
```

Outliers

```
# Assessing Outliers
```

```
outlierTest(fit) # Bonferonni p-value for most extreme obs
```

```
##
```

```
## No Studentized residuals with Bonferonni  $p < 0.05$ 
```


```
## Largest |rstudent|:
```

```
##               rstudent unadjusted p-value Bonferonni p
```

```
## Toyota Corolla  2.51597             0.01838       0.58816
```

```
qqPlot(fit, main="QQ Plot") #qq plot for studentized resid
```

QQ Plot




Influential Observations

```
# added variable plots  
###av.Plots(fit)
```


Cook's D plot


```
# identify D values > 4/(n-k-1)
cutoff <- 4/((nrow(mtcars)-length(fit$coefficients)-2))
plot(fit, which=4, cook.levels=cutoff)
```



Non-normality

```
# Normality of Residuals  
# qq plot for studentized resid  
qqPlot(fit, main="QQ Plot")
```

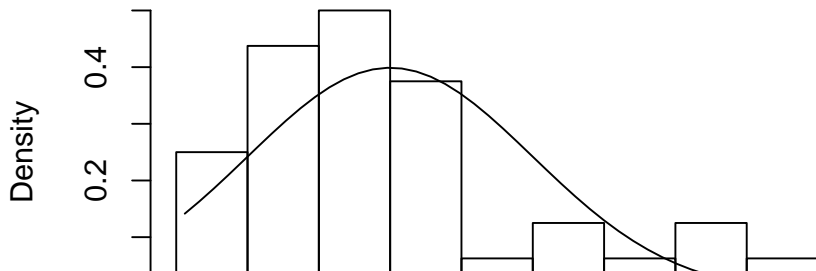
QQ Plot



distribution of studentized residuals

```
library(MASS)
sresid <- studres(fit)
hist(sresid, freq=FALSE,
     main="Distribution of Studentized Residuals")
xfit<-seq(min(sresid),max(sresid),length=40)
yfit<-dnorm(xfit)
lines(xfit, yfit)
```

Distribution of Studentized Residuals



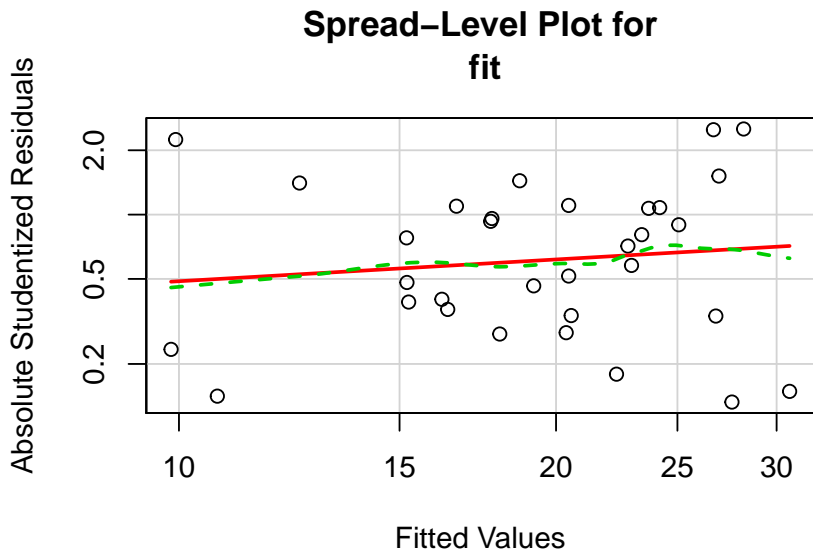
Non-constant Error Variance

```
# Evaluate homoscedasticity  
# non-constant error variance test  
ncvTest(fit)
```

```
## Non-constant Variance Score Test  
## Variance formula: ~ fitted.values  
## Chisquare = 1.429672    Df = 1    p = 0.231818
```

plot studentized residuals vs. fitted values

```
spreadLevelPlot(fit)
```



Multi-collinearity

```
# Evaluate Collinearity  
vif(fit) # variance inflation factors
```

```
##      disp      hp      wt      drat  
## 8.209402 2.894373 5.096601 2.279547
```

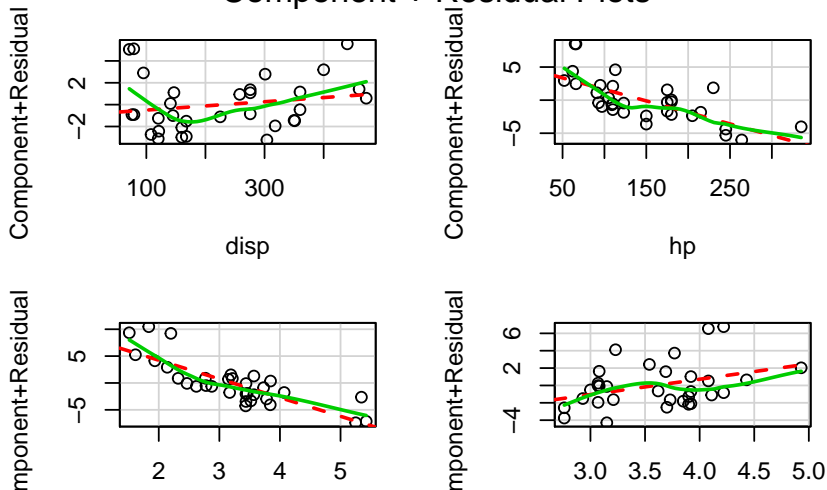
```
sqrt(vif(fit)) > 2
```

```
##  disp    hp    wt  drat  
## TRUE FALSE TRUE FALSE
```

Nonlinearity

```
# Evaluate Nonlinearity  
# component + residual plot  
crPlots(fit)
```

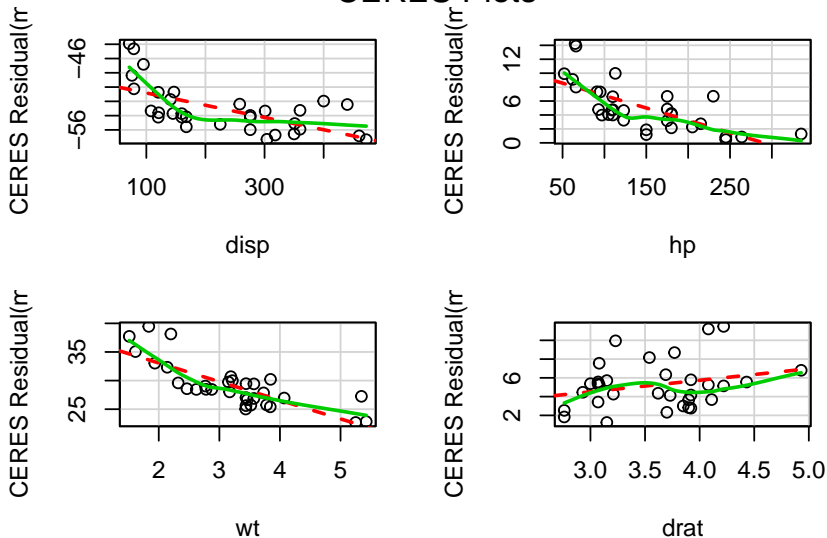
Component + Residual Plots



Ceres plots

```
ceresPlots(fit)
```

CERES Plots



Non-independence of Errors

```
# Test for Autocorrelated Errors  
durbinWatsonTest(fit)
```

```
## lag Autocorrelation D-W Statistic p-value  
## 1 0.100862 1.735915 0.284  
## Alternative hypothesis: rho != 0
```

Basic Graphs


Basic Graphs:

- Creating a Graph
- Density Plots
- Dot Plots
- Bar Plots
- Boxplots
- Scatter Plots

Creating a Graph

```
with(mtcars, plot(wt, mpg) )  
abline(lm(mpg~wt))  
title("Regression of MPG on Weight")
```

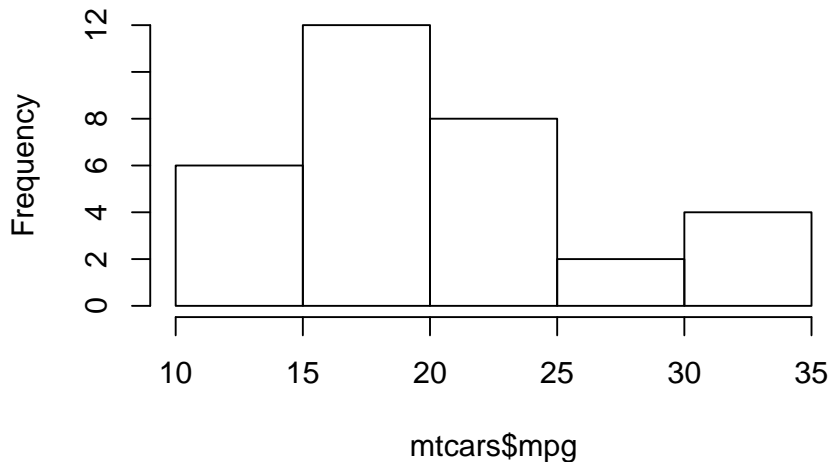
Regression of MPG on Weight



Density Plots

```
hist(mtcars$mpg)
```

Histogram of mtcars\$mpg

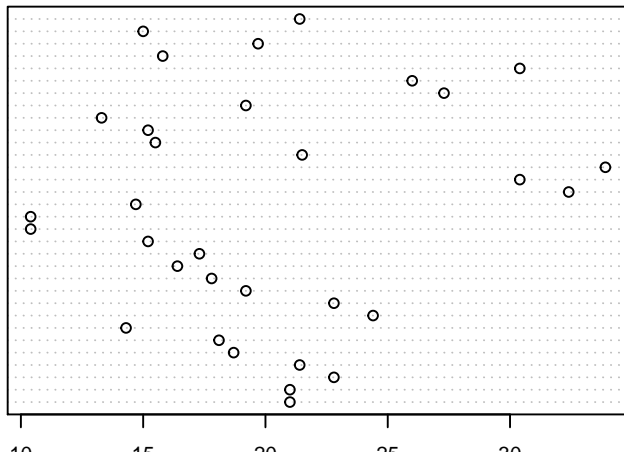


Dot Plots

```
dotchart(mtcars$mpg, labels=row.names(mtcars), cex=.7,  
         main="Gas Milage for Car Models",  
         xlab="Miles Per Gallon")
```


Gas Milage for Car Models

Volvo 142E
Maserati Bora
Ferrari Dino
Ford Pantera L
Lotus Europa
Porsche 914-2
Fiat X1-9
Pontiac Firebird
Camaro Z28
AMC Javelin
Dodge Challenger
Toyota Corona
Toyota Corolla
Honda Civic
Fiat 128
Chrysler Imperial
Lincoln Continental
Cadillac Fleetwood
Merc 450SLC
Merc 450SL
Merc 450SE
Merc 280C
Merc 280
Merc 230
Merc 240D
Duster 360
Valiant
Hornet Sportabout
Hornet 4 Drive
Datsun 710
Mazda RX4 Wag
Mazda RX4



Bar Plots

```
counts <- table(mtcars$gear)
barplot(counts, main="Car Distribution",
        xlab="Number of Gears")
```

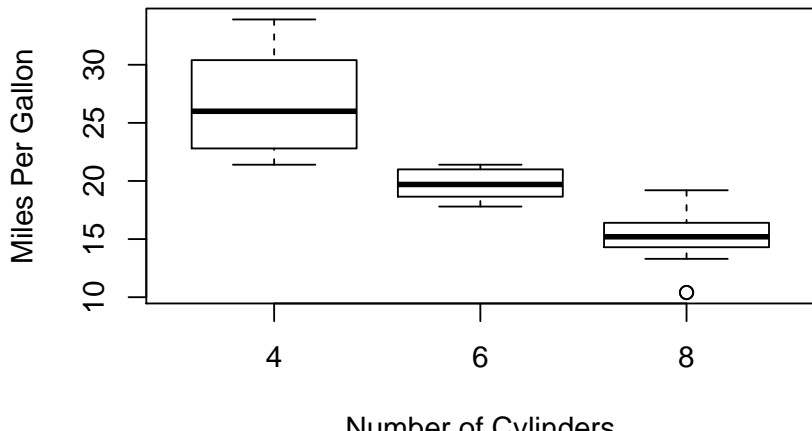


Boxplots

```
# Boxplot of MPG by Car Cylinders
```

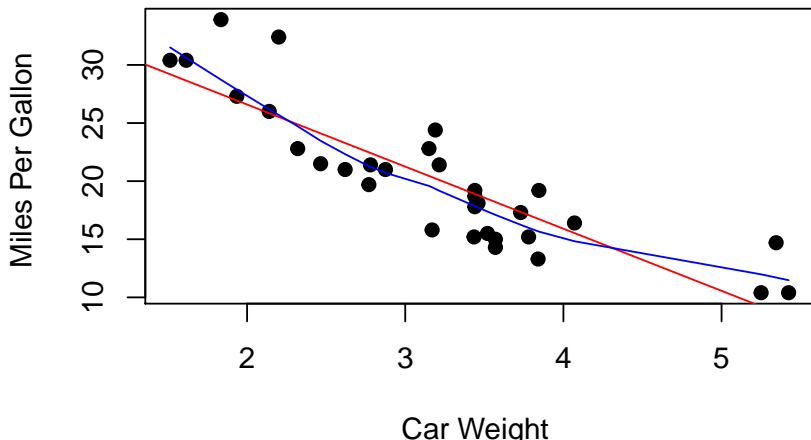
```
boxplot(mpg~cyl,data=mtcars, main="Car Milage Data",  
        xlab="Number of Cylinders", ylab="Miles Per Gallon")
```

Car Milage Data



Scatter Plots

```
plot(wt, mpg,  
     xlab="Car Weight ", ylab="Miles Per Gallon ", pch=19)  
#Add fit lines  
abline(lm(mpg~wt), col="red") # regression line (y~x)  
lines(lowess(wt,mpg), col="blue") # lowess line (x,y)
```



Scatterplot Matrices

```
pairs(~mpg+disp+drat+wt,data=mtcars)
```

